

Pedro Torres (pt397) and Ivan Zheng (iz60)

Project 3 Writeup

Code Implementation

`get_avail_ino`

We implemented it by first reading the inode bitmap from the disk. Once we have the bitmap, we iterated through from inode 0 to the max inode number and searched for the first bit that was set to 0. We returned this inode number

`get_avail_blkno`

We implemented it by first reading the block bitmap from the disk. Once we have the bitmap, we iterated through from block 0 to the max block number and searched for the first bit that was set to 0. We returned this block number + the start position of data blocks in the drive.

`readi`

We first found the block the inode belonged to ($\text{ino} / \text{INO_FACTOR}$). The `INO_FACTOR` was derived by $\text{sizeof}(\text{inode}) / \text{BLOCK_SIZE}$. We also found the inode offset in the block. We used `bio_read` to grab the inode from the disk, then stored the inode in the pointer given to us from the parameters of the function.

`writei`

Similar to `readi`, we first found the block and offset the inode belonged to using the `INO_FACTOR`. We called `bio_read` to get the block data, grabbed the inode pointer from the offset, copied the new inode to the pointer, then wrote back to the block.

`dir_find`

To implement `dir_find`, we first called `readi` to get the inode of the current directory. We then get the first data block that the inode points to. We then iterate through all the data blocks pointed to by the directory inode and for each data block, we iterate through all the dirents to check if the target file exists. We copy this dirent to the dirent given to us in the function parameters.

`dir_add`

To implement `dir_add`, we first call `dir_find` to see if the file we are trying to add already exists. We then check if the parent directory has space to accommodate a new dirent. If it does not, we add a new data

block. We then iterate through the directory's data blocks and find a dirent that is not populated. We fill out this dirent for the new directory we are adding.

`get_node_by_path`

To implement `get_node_by_path`, we first iteratively loop through the path, calling `dir_find` at each level and returning once we finish, we save the inode from path to function parameter inode pointer and return 0, and a negative error otherwise. Each level is separated by a slash in the path string.

`rufs_mkfs`

In implementing `rufs_mkfs`, initiate the disk with the `diskfile_path`, and initiate all data structures we need for rufs. This includes the superblock, inode and data bitmaps, and the root inode for the root path `"/`". We write any relevant changes back into the disk.

`rufs_init`

In implementing `rufs_init`, we allocate the memory for the bitmaps and superblocks, and we read from memory to fetch data needed to populate the bitmaps and superblocks.

`rufs_destroy`

We free the malloced bitmaps and superblocks and call `dev_close`.

`rufs_getattr`

We call `get_node_by_path` to get the inode from the path, then we save the inode's `vstat` member into `stbuf`.

`rufs_opendir`

We call `get_node_by_path` to get the inode from the path, then return 0 if it exists, other returns -1.

`rufs_readdir`

In implementing `rufs_readdir`, we call `get_node_by_path` to get the directory inode, We then iterate through the inodes data blocks and therefore its dirents to call the filler function to populate with the directories file entries.

`rufs_mkdir`

To implement `rufs_mkdir`, we call `get_node_by_path` to get the parent directory from path. We then get the next available inode to create the new directory's inode. We then save this inode into the datablock of the parent directory using `dir_add`. We use `writel` to write the new inode back into the block.

`rufs_create`

To implement `rufs_create`, we call `get_node_by_path` to get the inode of the parent directory. We then get the next available inode to create the new file's inode. We then save this inode into the data block of the parent directory using `dir_add`. We use `writel` to write the new inode back into the block.

`rufs_open`

To implement `rufs_open`, we call `get_node_by_path` and if it exists we return 0, else -1.

`rufs_read`

To implement `rufs_read`, we call `get_node_by_path` to get the file's inode. We then calculate the start and end block from the offset and size requested. We then iterate through and use `memcpy` to copy from the inode's data blocks into the buffer we are given.

`rufs_write`

To implement `rufs_write`, we call `get_node_by_path` to get the file's inode. We then calculate the start and end block from the offset and size requested. If we need to add more data blocks to accommodate, we do. We then iterate through and use `memcpy` to copy from the buffer to the inode's data blocks.